

CS3401 : Operating Systems

Project 2:

Loading Kernel Modules and Listing Tasks

Mohamed Eliskandrani

900140998

Introduction:

There are two parts to this project. The aim of the first is to load kernel modules into the linux kernel dynamically and the aim of the second was to traverse and display all tasks in the system both linearly and using dfs. There are several states that the task can take which can be shown by the project. The first and most common of these is the running state which indicates that the process is active and serving its requests. Sleeping means the process is awaiting the resources to run. Interruptible sleep means the process is waiting for a particular time slot or event to occur. Uninterruptible sleep means that it will awake when waited-upon resources become available to it. Stopped means the process has ended or has been terminated but has not been removed from the process table to ensure that parent processes know that the child has terminated successfully. Zombie is a state where the parent dies before it releases its child in the process table.

Code:

The code uses four linux libraries `init.h`, `kernel.h`, `module.h` and `sched.h`. `Init.h` is the library that contains the macros needed for the project. `Module.h` is needed by all kernel modules. `Kernel.h` is needed to access the `printk` function under `KERN_ALERT`. `Sched.h` is needed to access the functions `list_for_each`, `list_entry` and `for_each_process`. `For_each_process` iterates over all `task_struct` variables defined in `sched.h`, `list_for_each` iterates over the linked list of tasks and `list_entry` gets the `task_struct` for the current entry. The dfs function works by using the `list_for_each` and `list_entry` functions to recursively iterate over the list by calling `dfs(child)` for each child which in turn gets called for child's children and only when the functions have exhausted all of the children of a parent task does it move on to the next task and so on. You can check that the dfs is working by checking the parent id of the tasks listed.

Difficulty and Complications:

Most of the time spent on the project was used towards researching how the aforementioned functions work and how to use them to traverse all tasks in a dfs manner and researching how the linux kernel is structured and how tasks can be accessed using the `sched.h` library. If I were to do this project again I would probably do a lot more research before attempting to write code and wasting time on code that is already implemented using functions and macros.